

Parallel Matrix Multiplication

Parallel Matrix Multiplication

- Computing $C=A*B$
- Using basic algorithms
 - N^3
- Variables
 - Data layout
 - Topology of machine
 - Scheduling communications

Multiple Algorithms

- Divide the computation and data on the rows of the product matrix C
 - Each process holds n/p rows of A , entire B , and calculate n/p rows of C
 - Communications?
 - Scalable?

SUMMA Algorithm

- SUMMA = Scalable Universal Matrix Multiply
- Simple and easy to generalize
- Presentation from van de Geijn and Watts
 - www.netlib.org/lapack/lawns/lawn96.ps
 - Similar ideas appeared many times
- Used in practice in PBLAS = Parallel BLAS
 - www.netlib.org/lapack/lawns/lawn100.ps

Standard Algorithm

- standard matrix multiplication is computed using a sequence of inner product computations
- Assuming all $C_{i,j}$ are initialized to 0, the outer-product is

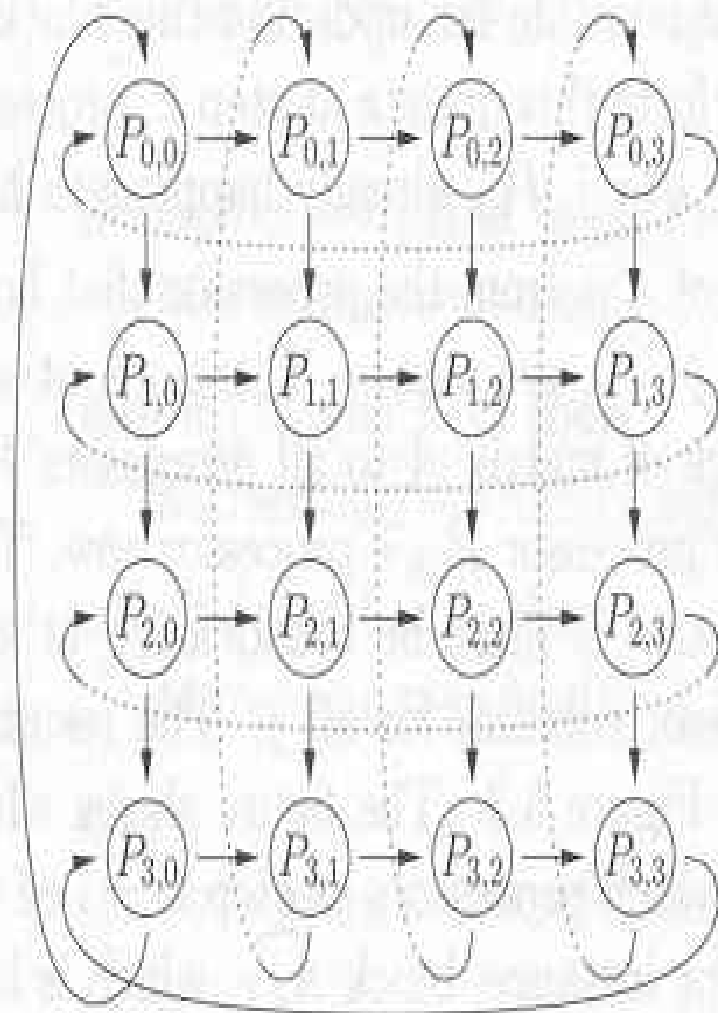
```
for i = 0 to n-1 do
  for j = 0 to n-1 do
  {
     $C_{i,j} = 0$ 
    for k = 0 to n-1 do
       $C_{i,j} = C_{i,j} + A_{i,k} \times B_{k,j}$ 
    }
  }
```

Outer-Product Algorithm

- Assuming all $C_{i,j}$ are initialized to 0, the outer-product is
 - for $k = 0$ to $n-1$ do
 - for $i = 0$ to $n-1$ do
 - for $j = 0$ to $n-1$ do
 - $C_{i,j} = C_{i,j} + A_{i,k} \times B_{k,j}$
- This outer-product leads to a simple and elegant parallelization on a torus of processors.
- At each step k , all $C_{i,j}$ are updated

Matrix Multiplication on a Grid


- Assume $p = q^2$ processors.
- Assume the matrix is $n \times n$ and that q divides n .
- The matrices are stored on the square $q \times q$ grid processors
- If $m = n/q$, each process holds a $m \times m$ block of each matrix.
- Technically, processor $P_{i,j}$ for $0 \leq i,j < q$ holds matrix blocks $A_{i,j}$, $B_{i,j}$, and $C_{i,j}$.
- This is illustrated on the next slide.



$\widehat{A}_{0,0}$	$\widehat{A}_{0,1}$	$\widehat{A}_{0,2}$	$\widehat{A}_{0,3}$
$\widehat{A}_{1,0}$	$\widehat{A}_{1,1}$	$\widehat{A}_{1,2}$	$\widehat{A}_{1,3}$
$\widehat{A}_{2,0}$	$\widehat{A}_{2,1}$	$\widehat{A}_{2,2}$	$\widehat{A}_{2,3}$
$\widehat{A}_{3,0}$	$\widehat{A}_{3,1}$	$\widehat{A}_{3,2}$	$\widehat{A}_{3,3}$

FIGURE 5.2: 2-D block distribution of an $n \times n$ matrix ($n = 24$) on a unidirectional grid/torus of p processors ($p = 4^2 = 16$).

SUMMA



For $k=0$ to $q-1$

for all $I = 1$ to p_r ... in parallel

owner of $A(I,k)$ broadcasts it to whole processor row

for all $J = 1$ to p_c ... in parallel

owner of $B(k,J)$ broadcasts it to whole processor column

Receive $A(I,k)$ into $Acol$

Receive $B(k,J)$ into $Brow$

$C(myrow, mycol) = C(myrow, mycol) + Acol * Brow$

Outer-Product Algorithm

- This algorithm can be summarized in terms of matrix **blocks** and matrix multiplications as

```
for  $k = 0$  to  $q - 1$  do
  for  $i = 0$  to  $q - 1$  do
    for  $j = 0$  to  $q - 1$  do
       $\widehat{C}_{i,j} \leftarrow \widehat{C}_{i,j} + \widehat{A}_{i,k} \times \widehat{B}_{k,j}$ 
```

- Next we consider executing this algorithm on a torus of $p = q^2$ processors.
- Processor $P_{i,j}$ holds block $C_{i,j}$ and updates it each step.
- To perform Step k , $P_{i,j}$ needs blocks $A_{i,k}$ & $B_{k,j}$.
- At Step $k=j$, $P_{i,j}$ already holds block $A_{i,j}$.
- For all other steps, $P_{i,j}$ must obtain $A_{i,k}$ from $P_{i,k}$.

Outer-Product Algorithm

- This is true for all processors $P_{i,j}$ with $j \neq k$.
- Note this means that at step k , processor $P_{i,k}$ must broadcast its block of matrix A to all processors $P_{i,j}$ on its row.
 - This is true for all rows i , as well.
- Similarly, blocks of matrix B must be broadcast at step k by $P_{k,j}$ to all processors on column— and for all j .
- The resulting communication pattern is shown on the next slide.



Matrix A Matrix B 

FIGURE 5.3: Communications of blocks of matrices A and B at step $k = 1$ of the outer-product matrix multiplication algorithm on a 4×4 torus of processors.

Algorithm 1 matmul function implementation if row and col communicators are used. Your program should have the exact structure. Pay close attention to how MPI_Bcast is called. You will need to implement a matmulAdd function.

Input: myrank, proc_grid_sz, block_sz, myA, myB, myC

Output: none

```
1: double **buffA, **buffB
2: buffA = alloc_2d_double(block_sz, block_sz)
3: buffB = alloc_2d_double(block_sz, block_sz)
4: create grid comm and get coordinates {Follow cartesian example}
5: create row comm
6: create col comm
7: for  $k \leftarrow 0$  to  $proc\_grid\_sz - 1$  do
8:   if coordinates[1] =  $k$  then
9:     copy items of myA to buffA
10:  end if
11:  MPI_Bcast(*buffA, block_sz*block_sz, MPI_DOUBLE, k, row_comm){*buffA specifies the
    starting memory location of the matrix buffA.}
```

```
12:  if coordinates[0] = k then
13:      copy items of myB to buffB
14:  end if
15:  MPI_Bcast(*buffB, block_sz*block_sz, MPI_DOUBLE, k, col_comm)
16:  if coordinates[0] = k && coordinates[1] = k then
17:      matmulAdd(myC, myA, myB, block_sz)
18:  else if coordinates[0] = k then
19:      matmulAdd(myC, buffA, myB, block_sz)
20:  else if coordinates[1] = k then
21:      matmulAdd(myC, myA, buffB, block_sz)
22:  else
23:      matmulAdd(myC, buffA, buffB, block_sz)
24:  end if
25: end for
```

Outer Product Algorithm Steps

- Statement 1 declares the square blocks of the three matrices stored by each processor.
 - The matrix C is assumed to be initialized to zero
 - Arrays A & B contain sub-matrices in PEs in Fig 5.2
- Statements 2&3 declare two helper buffers used by PEs
- The q steps of program occur in lines 7-25
- In statements 8-11, all q processors in column k broadcast (in parallel) their block of A to the processors in each of their rows.
- Statements 12-15 implement similar broadcasts of blocks of matrix B along processor columns.

Outer Product Algorithm Steps (cont)

- Comments:
 - When preceding broadcasts are complete, each PE holds all the needed blocks.
 - Each processor will multiply a block of A by a block of B and adds the result to the block of C, for which it is responsible.
 - The algorithm uses the notation `MatrixMultiplyAdd()` for PE matrix block operations of $C_{i,j} \leftarrow C_{i,j} + A_{i,k} B_{k,j}$.
- In lines 16-17, if the PE is on both row k & column k, then it can just multiply the two blocks of A and B that it holds.
- Lines 18-19: If the PE is on row k but not on column k, then it will multiply the block of A that it receives with the block of B that it holds.

Outer Product Algorithm Steps (cont)

- Lines 20-21: Similarly, if a PE is on column k but not row k , then it multiplies the block of A it holds with the block of B it just received.
- Lines 22-23 (General Case): If a PE is neither on row k or column k , then it will multiply the block of A it receives with the block of B that it receives.

Generalization of Matrix Multiply:

- By allotting rectangular blocks of Matrix A and B to processors, the preceding algorithm can be adapted to work for non-square matrix products.